

Analyzing and Preprocessing Data

January 13, 2026

##Pre-processing Data In this project, we will explore several methods of data preprocessing for Machine Learning.

Problem: Ensuring proper data cleaning and preprocessing to prepare datasets for Machine Learning applications.

```
[ ]: #Importing python libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Here, we're going to use the E-commerce Customer Behavior dataset from Kaggle. Courtesy: <https://www.kaggle.com/dhairyajeetsingh/ecommerce-customer-behavior-dataset>

```
[ ]: # Reading Dataset from local csv file
df_cust = pd.read_csv('ecommerce_customer_churn_dataset.csv')
```

Note: Pandas can read datasets in various formats:

CSV - read_csv(<filename>) : Reads comma separated files

TSV - read_csv(<filename>, sep= "") : Reads a "tab" delimited file, you can replace separator with :, | and read different types of delimited file

Excel - read_excel(<filename>, sheet_name=<worksheet_name>) : Reads .xls or .xlsx file

Let us check dimension of Dataset

```
[ ]: print(df_cust.shape)
```

(50000, 25)

This data set has 25 features and 50000 samples

Let's examine few rows from this dataset

```
[ ]: print(df_cust.head())
```

	Age	Gender	Country	City	Membership_Years	Login_Frequency	\
0	43.0	Male	France	Marseille	2.9	14.0	
1	36.0	Male	UK	Manchester	1.6	15.0	
2	45.0	Female	Canada	Vancouver	2.9	10.0	
3	56.0	Female	USA	New York	2.6	10.0	

```

4 35.0    Male   India      Delhi        3.1        29.0
          Session_Duration_Avg  Pages_Per_Session  Cart_Abandonment_Rate \
0                  27.4            6.0                50.6
1                  42.7            10.3               37.7
2                  24.8             1.6                70.9
3                  38.4            14.8               41.7
4                  51.4            NaN               19.1

          Wishlist_Items  Total_Purchases  Average_Order_Value \
0                  3.0            9.0                94.72
1                  1.0            19.5               82.45
2                  1.0            9.1                165.52
3                  9.0            15.0               147.33
4                  9.0            32.5               141.30

          Days_Since_Last_Purchase  Discount_Usage_Rate  Returns_Rate \
0                      34.0            46.40              2.0
1                      71.0            57.96              9.2
2                      11.0            12.24             11.5
3                      47.0            44.10              5.4
4                      73.0            25.20              5.5

          Email_Open_Rate  Customer_Service_Calls  Product_Reviews_Written \
0                  17.9            9.0                4.0
1                  42.8            7.0                3.0
2                  0.0            4.0                1.0
3                  41.4            2.0                5.0
4                  37.9            1.0               11.0

          Social_Media_Engagement_Score  Mobile_App_Usage  Payment_Method_Diversity \
0                      16.3            20.8                1.0
1                      NaN            23.3                3.0
2                      NaN             8.8                NaN
3                      85.9            31.0                3.0
4                      83.0            50.4                4.0

          Lifetime_Value  Credit_Balance  Churned Signup_Quarter
0          953.33        2278.0       0        Q1
1         1067.47        3028.0       0        Q4
2         1289.75        2317.0       0        Q4
3         2340.92        2674.0       0        Q1
4         3041.29        5354.0       0        Q4

```

By default, the head() function displays the first 5 rows and up to 20 columns. To view additional rows and columns, these defaults can be adjusted, allowing the DataFrame to be formatted into a more readable table.

```
[ ]: pd.set_option('display.max_columns', 25)
      print(df_cust.head(3).to_string())
```

Age	Gender	Country	City	Membership_Years	Login_Frequency
Session_Duration_Avg	Pages_Per_Session	Cart_Abandonment_Rate	Wishlist_Items		
Total_Purchases	Average_Order_Value	Days_Since_Last_Purchase			
Discount_Usage_Rate	Returns_Rate	Email_Open_Rate	Customer_Service_Calls		
Product_Reviews_Written	Social_Media_Engagement_Score	Mobile_App_Usage			
Payment_Method_Diversity	Lifetime_Value	Credit_Balance	Churned	Signup_Quarter	
0	43.0	Male	France	Marseille	2.9
27.4		6.0		50.6	3.0
94.72			34.0	46.40	2.0
17.9			9.0	4.0	
16.3		20.8		1.0	953.33
0		Q1			2278.0
1	36.0	Male	UK	Manchester	1.6
42.7		10.3		37.7	1.0
82.45			71.0	57.96	9.2
42.8			7.0	3.0	
NaN		23.3		3.0	1067.47
0		Q4			3028.0
2	45.0	Female	Canada	Vancouver	2.9
24.8		1.6		70.9	1.0
165.52			11.0	12.24	11.5
0.0			4.0	1.0	
NaN		8.8		NaN	1289.75
0		Q4			2317.0

Let us see the information about dataset

`info()` - gives us a quick snapshot of DataFrame. It shows us how many rows, columns & its data type, and whether there are any missing values

```
[ ]: print(df_cust.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 25 columns):
 #   Column           Non-Null Count   Dtype  
--- 
 0   Age              47505 non-null    float64
 1   Gender            50000 non-null    object 
 2   Country           50000 non-null    object 
 3   City              50000 non-null    object 
 4   Membership_Years  50000 non-null    float64
 5   Login_Frequency   50000 non-null    float64
 6   Session_Duration_Avg  46601 non-null    float64
 7   Pages_Per_Session 47000 non-null    float64
 8   Cart_Abandonment_Rate 50000 non-null    float64
```

```

9   Wishlist_Items           46000 non-null float64
10  Total_Purchases          50000 non-null float64
11  Average_Order_Value      50000 non-null float64
12  Days_Since_Last_Purchase 47000 non-null float64
13  Discount_Usage_Rate      46500 non-null float64
14  Returns_Rate              45509 non-null float64
15  Email_Open_Rate           47472 non-null float64
16  Customer_Service_Calls    49832 non-null float64
17  Product_Reviews_Written    46500 non-null float64
18  Social_Media_Engagement_Score 44000 non-null float64
19  Mobile_App_Usage          45000 non-null float64
20  Payment_Method_Diversity    47500 non-null float64
21  Lifetime_Value             50000 non-null float64
22  Credit_Balance             44500 non-null float64
23  Churned                     50000 non-null int64
24  Signup_Quarter              50000 non-null object
dtypes: float64(20), int64(1), object(4)
memory usage: 9.5+ MB
None

```

Looking at the dataset, you'll notice a mix of float, integer, and text/object columns. Some columns also have missing entries, which shows up when the non-null count is below 50,000.

```
[ ]: df_cust.columns[df_cust.isna().any()].tolist()
```

```
[ ]: ['Age',
       'Session_Duration_Avg',
       'Pages_Per_Session',
       'Wishlist_Items',
       'Days_Since_Last_Purchase',
       'Discount_Usage_Rate',
       'Returns_Rate',
       'Email_Open_Rate',
       'Customer_Service_Calls',
       'Product_Reviews_Written',
       'Social_Media_Engagement_Score',
       'Mobile_App_Usage',
       'Payment_Method_Diversity',
       'Credit_Balance']
```

To clean things up, let's replace the nulls in Session_Duration_Avg, Returns_Rate, and Email_Open_Rate with their average values.

```
[ ]: avg_features = ['Session_Duration_Avg', 'Returns_Rate', 'Email_Open_Rate', ↴
                   'Discount_Usage_Rate', 'Payment_Method_Diversity', 'Credit_Balance', ↴
                   'Mobile_App_Usage']
```

`describe()` : generates summary statistics for numeric columns, like mean, median, min, max, and quartiles. It's a quick way to understand the distribution and spread of your dataset's values.

```
[ ]: df_cust[avg_features].describe()
```

```
[ ]: Session_Duration_Avg Returns_Rate Email_Open_Rate \
count      46601.000000 45509.000000 47472.000000
mean       27.660754    6.680913    20.937980
std        10.871013   6.143027    14.252561
min        1.000000    0.000000    0.000000
25%       19.700000   2.900000    9.900000
50%       26.800000   5.400000   19.700000
75%       34.700000   9.100000   30.400000
max       75.600000  99.615734  91.700000

Discount_Usage_Rate Payment_Method_Diversity Credit_Balance \
count      46500.000000 47500.000000 44500.000000
mean       41.997485    2.353874    1966.233258
std        21.373642   1.110012    1225.072166
min        0.240000    1.000000    0.000000
25%       25.300000    2.000000    1049.000000
50%       40.200000    2.000000    1896.000000
75%       57.000000    3.000000    2791.000000
max       116.640000   5.000000    7197.000000

Mobile_App_Usage
count      45000.000000
mean       19.371607
std        9.419252
min        0.000000
25%       12.500000
50%       18.600000
75%       25.500000
max       61.900000
```

Let's take a look at how the average feature values have changed in the data after we applied the transformation.

```
[ ]: from sklearn.impute import SimpleImputer
      from sklearn.compose import ColumnTransformer

      avgimputer = SimpleImputer(strategy='mean')
      col_trans = ColumnTransformer([('avg_imputer', avgimputer, avg_features)], remainder='passthrough').set_output(transform="pandas")
      df_cust_trans = col_trans.fit_transform(df_cust)
```

Here, We've imported SimpleImputer from the scikit-learn library.

Scikit-learn is a free and open-source Python package offering efficient tools for machine learning and data analysis.

SimpleImputer serves as a preprocessing utility that helps us handle missing numerical values in a

dataset.

```
[ ]: df_cust_trans.describe()
```

```
[ ]: avg_imputer__Session_Duration_Avg avg_imputer__Returns_Rate \
count 50000.000000 50000.000000
mean 27.660754 6.680913
std 10.494997 5.860648
min 1.000000 0.000000
25% 20.200000 3.100000
50% 27.660754 6.000000
75% 34.000000 8.600000
max 75.600000 99.615734

avg_imputer__Email_Open_Rate avg_imputer__Discount_Usage_Rate \
count 50000.000000 50000.000000
mean 20.937980 41.997485
std 13.887576 20.611979
min 0.000000 0.240000
25% 10.500000 26.500000
50% 20.800000 41.997485
75% 29.700000 55.560000
max 91.700000 116.640000

avg_imputer__Payment_Method_Diversity avg_imputer__Credit_Balance \
count 50000.000000 50000.000000
mean 2.353874 1966.233258
std 1.081905 1155.729342
min 1.000000 0.000000
25% 2.000000 1164.000000
50% 2.000000 1966.233258
75% 3.000000 2664.000000
max 5.000000 7197.000000

avg_imputer__Mobile_App_Usage remainder__Age \
count 50000.000000 47505.000000
mean 19.371607 37.802968
std 8.935878 11.834668
min 0.000000 5.000000
25% 13.200000 29.000000
50% 19.371607 38.000000
75% 24.600000 46.000000
max 61.900000 200.000000

remainder__Membership_Years remainder__Login_Frequency \
count 50000.000000 50000.000000
mean 2.984009 11.624660
std 2.059105 7.810657
```

min	0.100000	0.000000
25%	1.400000	6.000000
50%	2.500000	11.000000
75%	4.000000	17.000000
max	10.000000	46.000000
remainder__Pages_Per_Session	47000.000000	50000.000000
count	8.737811	57.079973
mean	3.778220	16.282723
std	1.000000	0.000000
min	6.000000	46.400000
25%	8.400000	58.100000
50%	11.200000	68.700000
75%	24.100000	143.743350
remainder__Wishlist_Items	46000.000000	50000.000000
count	4.298391	13.111576
mean	3.189754	7.017312
std	0.000000	-13.000000
min	2.000000	8.000000
25%	4.000000	12.000000
50%	6.000000	17.000000
75%	28.000000	128.700000
remainder__Average_Order_Value	50000.000000	47000.000000
count	123.117330	29.792872
mean	175.569714	29.695062
std	26.380000	0.000000
min	87.050000	9.000000
25%	112.970000	21.000000
50%	144.440000	41.000000
75%	9666.379178	287.000000
remainder__Customer_Service_Calls	49832.000000	46500.000000
count	5.681831	2.853312
mean	2.676052	2.328948
std	0.000000	0.000000
min	4.000000	1.000000
25%	5.000000	2.000000
50%	7.000000	4.000000
75%	21.000000	21.000000
remainder__Social_Media_Engagement_Score	remainder__Lifetime_Value	

count	44000.000000	50000.000000
mean	29.364466	1440.626292
std	20.574021	907.249443
min	0.000000	0.000000
25%	13.200000	789.817500
50%	27.600000	1243.415000
75%	43.100000	1874.000000
max	100.000000	8987.240000
remainder__Churned		
count	50000.000000	
mean	0.289000	
std	0.453302	
min	0.000000	
25%	0.000000	
50%	0.000000	
75%	1.000000	
max	1.000000	

Now let's use SimpleImputer to clean up the dataset by replacing nulls in Age, Wishlist_Items, and Customer_service_calls with the median values.

```
[ ]: med_features = ['Age', 'Wishlist_Items', 'Customer_Service_Calls']
medimputer = SimpleImputer(strategy='median')
col_trans = ColumnTransformer([('avg_imputer', avgimputer, avg_features),
                             ('med_imputer', medimputer, med_features)], remainder='passthrough').set_output(transform="pandas")
df_cust_trans = col_trans.fit_transform(df_cust)
```

ColumnTransformer() - lets you apply different preprocessing steps to different columns at once in a single pipeline.

```
[ ]: df_cust_trans.describe()
```

	avg_imputer__Session_Duration_Avg	avg_imputer__Returns_Rate \
count	50000.000000	50000.000000
mean	27.660754	6.680913
std	10.494997	5.860648
min	1.000000	0.000000
25%	20.200000	3.100000
50%	27.660754	6.000000
75%	34.000000	8.600000
max	75.600000	99.615734
	avg_imputer__Email_Open_Rate	avg_imputer__Discount_Usage_Rate \
count	50000.000000	50000.000000
mean	20.937980	41.997485
std	13.887576	20.611979

min	0.000000	0.240000
25%	10.500000	26.500000
50%	20.800000	41.997485
75%	29.700000	55.560000
max	91.700000	116.640000
avg_imputer__Payment_Method_Diversity	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	2.353874	1966.233258
std	1.081905	1155.729342
min	1.000000	0.000000
25%	2.000000	1164.000000
50%	2.000000	1966.233258
75%	3.000000	2664.000000
max	5.000000	7197.000000
avg_imputer__Mobile_App_Usage	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	19.371607	37.812800
std	8.935878	11.535688
min	0.000000	5.000000
25%	13.200000	30.000000
50%	19.371607	38.000000
75%	24.600000	45.000000
max	61.900000	200.000000
med_imputer__Wishlist_Items	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	4.274520	5.679540
std	3.060573	2.671844
min	0.000000	0.000000
25%	2.000000	4.000000
50%	4.000000	5.000000
75%	6.000000	7.000000
max	28.000000	21.000000
remainder__Membership_Years	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	2.984009	11.624660
std	2.059105	7.810657
min	0.100000	0.000000
25%	1.400000	6.000000
50%	2.500000	11.000000
75%	4.000000	17.000000
max	10.000000	46.000000
remainder__Pages_Per_Session	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	11.624660	7.810657
std	2.059105	0.100000
min	6.000000	11.000000
25%	17.000000	4.000000
50%	46.000000	2.500000
75%	10.000000	1.400000
max	50000.000000	0.100000
remainder__Cart_Abandonment_Rate	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	7.810657	11.624660
std	0.100000	2.059105
min	11.000000	6.000000
25%	4.000000	17.000000
50%	2.500000	10.000000
75%	1.400000	50000.000000
max	0.100000	46.000000

count	47000.000000	50000.000000
mean	8.737811	57.079973
std	3.778220	16.282723
min	1.000000	0.000000
25%	6.000000	46.400000
50%	8.400000	58.100000
75%	11.200000	68.700000
max	24.100000	143.743350
remainder__Total_Purchases	remainder__Average_Order_Value \	
count	50000.000000	50000.000000
mean	13.111576	123.117330
std	7.017312	175.569714
min	-13.000000	26.380000
25%	8.000000	87.050000
50%	12.000000	112.970000
75%	17.000000	144.440000
max	128.700000	9666.379178
remainder__Days_Since_Last_Purchase	\	
count	47000.000000	
mean	29.792872	
std	29.695062	
min	0.000000	
25%	9.000000	
50%	21.000000	
75%	41.000000	
max	287.000000	
remainder__Product_Reviews_Written	\	
count	46500.000000	
mean	2.853312	
std	2.328948	
min	0.000000	
25%	1.000000	
50%	2.000000	
75%	4.000000	
max	21.000000	
remainder__Social_Media_Engagement_Score	remainder__Lifetime_Value \	
count	44000.000000	50000.000000
mean	29.364466	1440.626292
std	20.574021	907.249443
min	0.000000	0.000000
25%	13.200000	789.817500
50%	27.600000	1243.415000
75%	43.100000	1874.000000

```

max           100.000000
8987.240000

      remainder__Churned
count      50000.000000
mean       0.289000
std        0.453302
min        0.000000
25%       0.000000
50%       0.000000
75%       1.000000
max        1.000000

```

Next, we'll apply SimpleImputer to fill missing values in selected columns using their most frequent entries.

```
[ ]: mode_features = [
    'Pages_Per_Session', 'Days_Since_Last_Purchase', 'Product_Reviews_Written', 'Social_Media_Engagement',
    modeimputer = SimpleImputer(strategy='most_frequent')
    col_trans = ColumnTransformer([('avg_imputer', avgimputer, avg_features),
                                    ('med_imputer', medimputer, med_features),
                                    ('mode_imputer', modeimputer, mode_features)],
                                 remainder='passthrough').
    set_output(transform="pandas")
df_cust_trans = col_trans.fit_transform(df_cust)
```

When you set remainder="passthrough", it just means any columns you didn't transform will still be carried forward instead of being dropped.

And if you use set_output(transform="pandas"), it will give you back a pandas DataFrame with proper column names, which makes the results much easier to read and work with.

```
[ ]: df_cust_trans.describe()
```

```

[ ]:      avg_imputer__Session_Duration_Avg  avg_imputer__Returns_Rate \
count          50000.000000                 50000.000000
mean         27.660754                  6.680913
std          10.494997                  5.860648
min         1.000000                  0.000000
25%        20.200000                  3.100000
50%        27.660754                  6.000000
75%        34.000000                  8.600000
max        75.600000                 99.615734

      avg_imputer__Email_Open_Rate  avg_imputer__Discount_Usage_Rate \
count          50000.000000                 50000.000000
mean         20.937980                  41.997485
std          13.887576                  20.611979
min         0.000000                  0.240000

```

25%	10.500000	26.500000
50%	20.800000	41.997485
75%	29.700000	55.560000
max	91.700000	116.640000
avg_imputer__Payment_Method_Diversity	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	2.353874	1966.233258
std	1.081905	1155.729342
min	1.000000	0.000000
25%	2.000000	1164.000000
50%	2.000000	1966.233258
75%	3.000000	2664.000000
max	5.000000	7197.000000
avg_imputer__Mobile_App_Usage	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	19.371607	37.812800
std	8.935878	11.535688
min	0.000000	5.000000
25%	13.200000	30.000000
50%	19.371607	38.000000
75%	24.600000	45.000000
max	61.900000	200.000000
med_imputer__Wishlist_Items	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	4.274520	5.679540
std	3.060573	2.671844
min	0.000000	0.000000
25%	2.000000	4.000000
50%	4.000000	5.000000
75%	6.000000	7.000000
max	28.000000	21.000000
mode_imputer__Pages_Per_Session	50000.000000	50000.000000
count	50000.000000	50000.000000
mean	8.705542	8.705542
std	3.665344	3.665344
min	1.000000	1.000000
25%	6.200000	6.200000
50%	8.200000	8.200000
75%	11.000000	11.000000
max	24.100000	24.100000
mode_imputer__Days_Since_Last_Purchase	50000.000000	50000.000000
count	50000.000000	50000.000000

```

mean                                28.065300
std                                 29.591317
min                                 0.000000
25%                                7.000000
50%                                19.000000
75%                                39.000000
max                                 287.000000

    mode_imputer__Product_Reviews_Written \
count                               50000.000000
mean                                2.723580
std                                 2.295194
min                                 0.000000
25%                                1.000000
50%                                2.000000
75%                                4.000000
max                                 21.000000

    mode_imputer__Social_Media_Engagement_Score \
count                               50000.000000
mean                                25.840730
std                                 21.530263
min                                 0.000000
25%                                5.700000
50%                                23.700000
75%                                40.700000
max                                 100.000000

    remainder__Membership_Years  remainder__Login_Frequency \
count                               50000.000000          50000.000000
mean                                2.984009          11.624660
std                                 2.059105          7.810657
min                                 0.100000          0.000000
25%                                1.400000          6.000000
50%                                2.500000         11.000000
75%                                4.000000         17.000000
max                                 10.000000        46.000000

    remainder__Cart_Abandonment_Rate  remainder__Total_Purchases \
count                               50000.000000          50000.000000
mean                                57.079973          13.111576
std                                 16.282723          7.017312
min                                 0.000000         -13.000000
25%                                46.400000          8.000000
50%                                58.100000         12.000000
75%                                68.700000         17.000000
max                                 143.743350        128.700000

```

```

        remainder__Average_Order_Value  remainder__Lifetime_Value \
count              50000.000000          50000.000000
mean             123.117330         1440.626292
std              175.569714         907.249443
min              26.380000          0.000000
25%             87.050000         789.817500
50%            112.970000        1243.415000
75%            144.440000        1874.000000
max             9666.379178        8987.240000

        remainder__Churned
count      50000.000000
mean       0.289000
std        0.453302
min       0.000000
25%       0.000000
50%       0.000000
75%       1.000000
max       1.000000

```

Let's quickly check if there are still any columns with missing data.

```
[ ]: df_cust_trans.columns[df_cust_trans.isna().any()].tolist()
```

```
[ ]: []
```

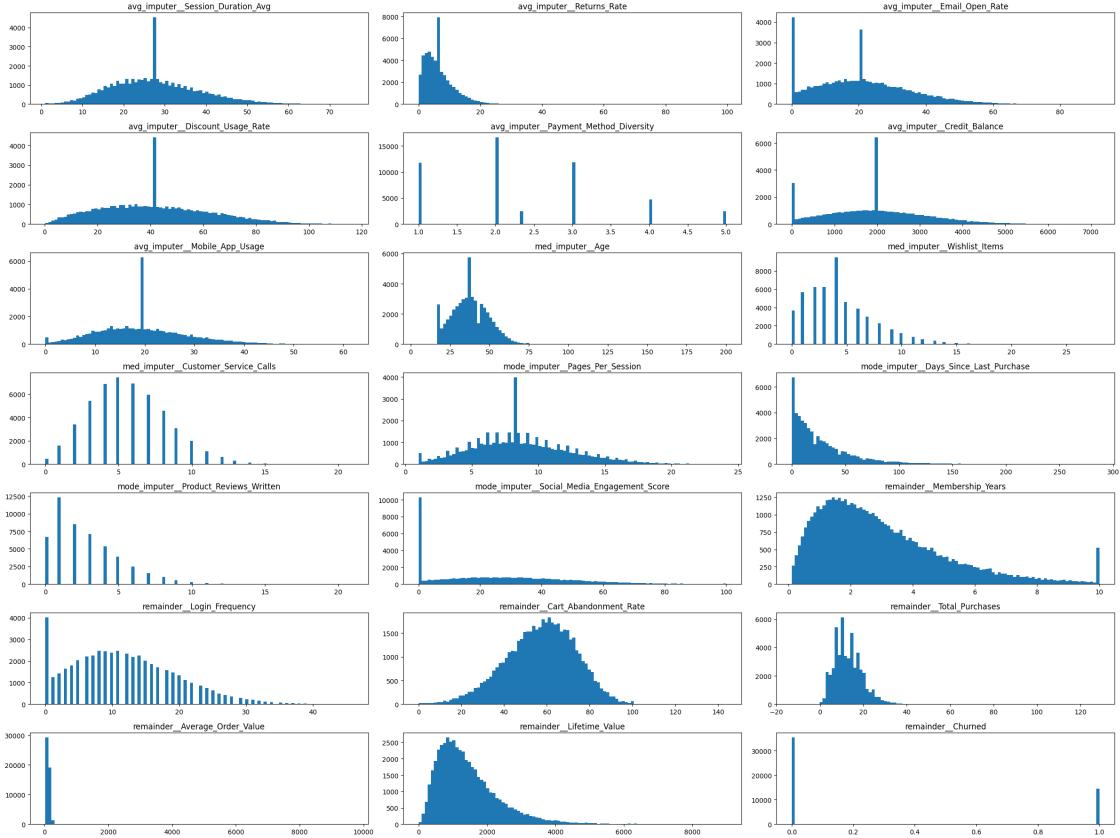
`isna()` - checks for missing values in a DataFrame or Series. It returns a boolean mask, with True where values are null and False otherwise.

`any()` - checks if at least one value in a Series or DataFrame is True.

`tolist()` - converts a pandas Series or DataFrame values into a plain Python list

Next, let's create visualizations to examine how the values are distributed across each column.

```
[ ]: numeric_cols=df_cust_trans.columns[df_cust_trans.dtypes != 'object']
fig, axes = plt.subplots(nrows=7,ncols=3, figsize=(24,18))
axes=axes.ravel()
idx=0
for col in numeric_cols:
    if idx < len(axes): # Added a check to prevent IndexError
        axes[idx].set_title(col)
        axes[idx].hist(df_cust_trans[col],bins=100)
    idx+=1
plt.tight_layout()
plt.show()
```

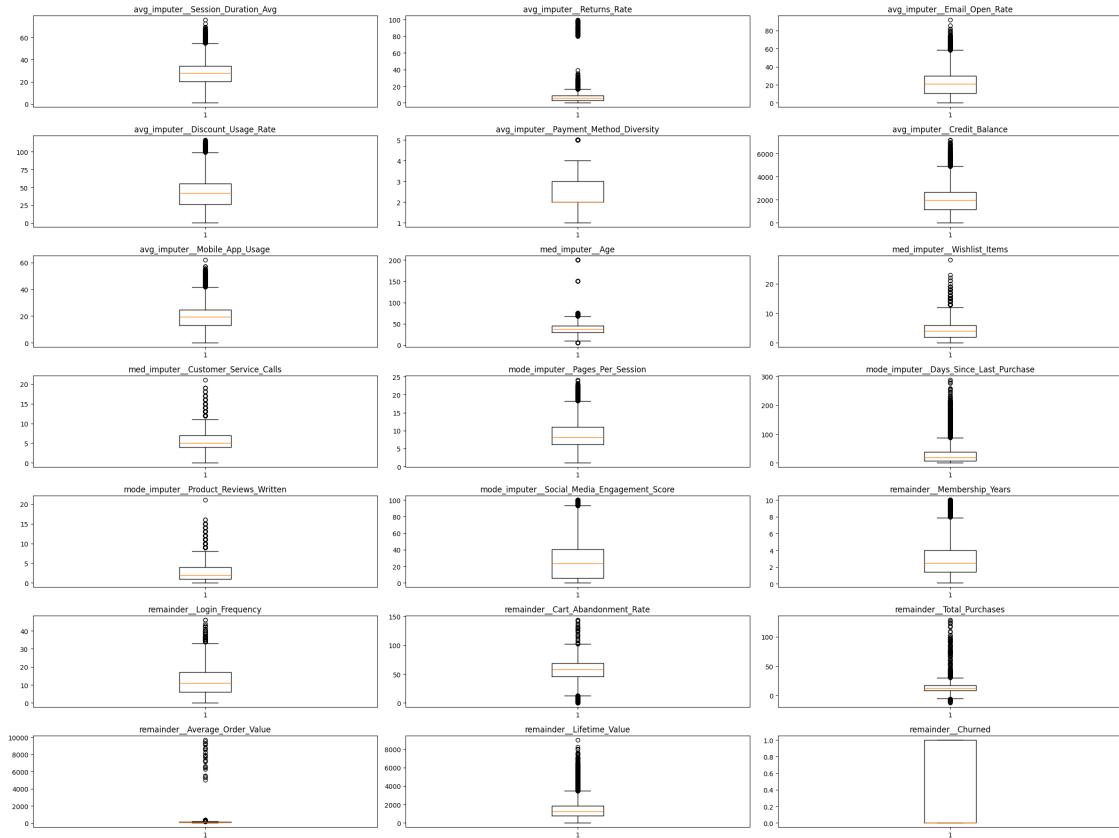


The above code builds a grid of histograms where each subplot shows the distribution of one numeric column. By arranging them side-by-side, thus you can easily see how features are spread out, spot skewness, or detect outliers.

- **Type of chart:** Histograms are used to show how values are distributed for each numeric column.
- **Grid:** A grid of plots is created (7×3 columns) to organize multiple charts on one canvas.
- **Subplot:** Each subplot corresponds to one column, with its own title and histogram.

The data isn't normalized — the ranges vary a lot and some columns have outliers. So, we'll start by spotting outliers with boxplots, and then normalize the data.

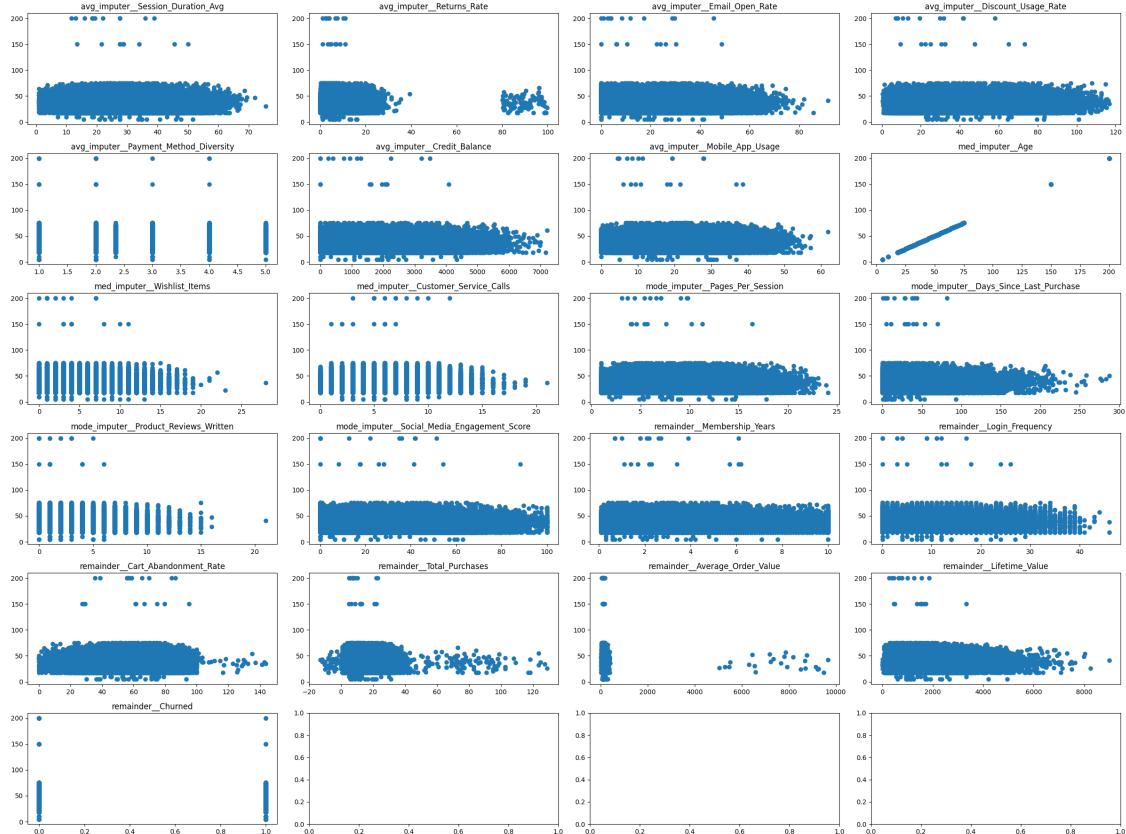
```
[ ]: fig, axes = plt.subplots(nrows=7, ncols=3, figsize=(24,18))
axes=axes.ravel()
idx=0
for col in numeric_cols:
    axes[idx].set_title(col)
    axes[idx].boxplot(df_cust_trans[col])
    idx+=1
plt.tight_layout()
plt.show()
```



[]:

Next, let's analyze scatter plots of the numeric columns in relation to the Age column.

```
[ ]: y=df_cust_trans['med_imputer__Age']
fig, axes = plt.subplots(nrows=6,ncols=4, figsize=(24,18))
axes=axes.ravel()
idx=0
for col in numeric_cols:
    axes[idx].set_title(col)
    axes[idx].scatter(df_cust_trans[col],y)
    idx+=1
plt.tight_layout()
plt.show()
```



Next, we'll handle the outliers by applying the Winsorize technique.

- Winsorizing is a statistical method used to limit extreme values in data.
- Instead of removing outliers completely, it replaces extreme values (both very high and very low) with the nearest values within a chosen percentile range.
- Example: If you Winsorize at the 5th and 95th percentiles, values below the 5th percentile are set to the 5th percentile value, and values above the 95th percentile.

```
[ ]: def winsorize(df, col, lower=0.01, upper=0.99):
    low_val = df[col].quantile(lower)
    high_val = df[col].quantile(upper)
    df[col] = np.clip(df[col], low_val, high_val)
    return df
```

```
[ ]: for col in numeric_cols:
    df_cust_trans = winsorize(df_cust_trans, col)
```

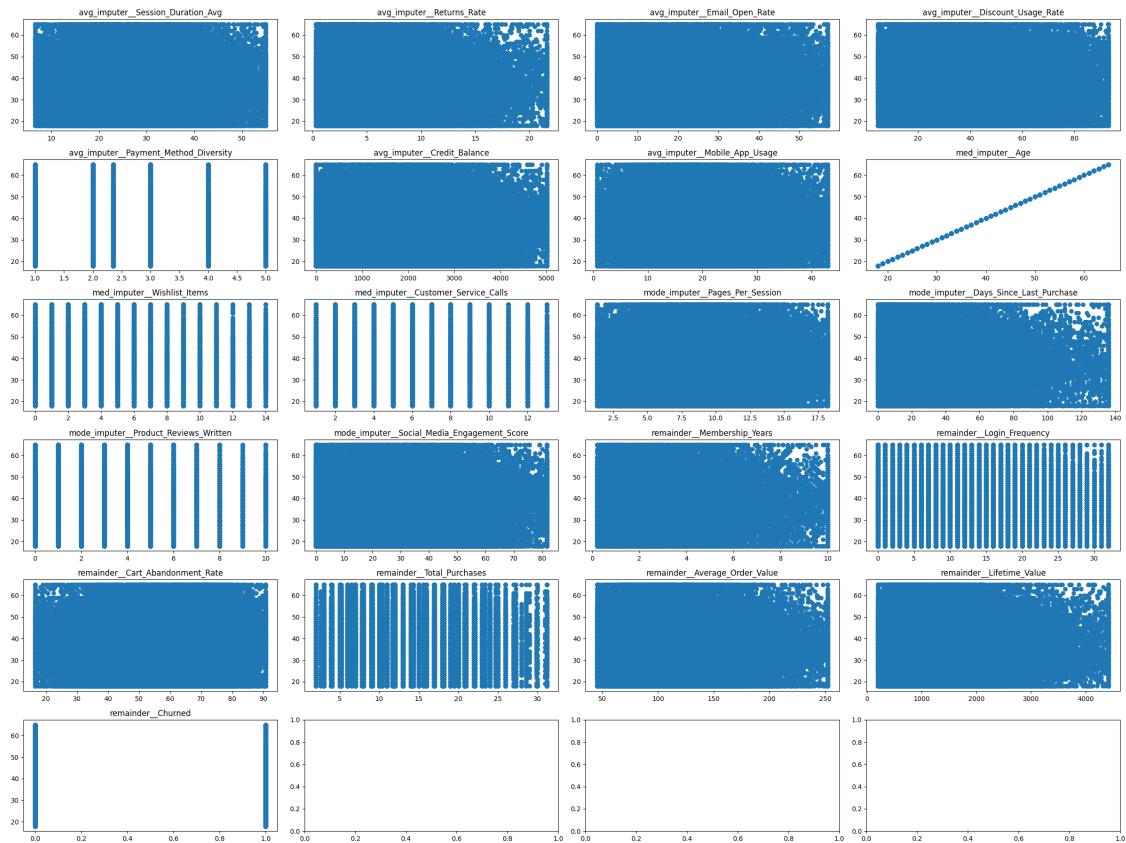
Now, we'll re-plot the scatter plots after removing the outliers to examine the cleaned data.

```
[ ]: y=df_cust_trans['med_imputer_Age']
fig, axes = plt.subplots(nrows=6, ncols=4, figsize=(24,18))
axes=axes.ravel()
```

```

idx=0
for col in numeric_cols:
    axes[idx].set_title(col)
    axes[idx].scatter(df_cust_trans[col],y)
    idx+=1
plt.tight_layout()
plt.show()

```



Let's scale the numeric columns so they're centered at 0 with variance 1, giving all features the same influence in training. We'll use StandardScaler to do this.

```

[ ]: from sklearn.preprocessing import StandardScaler

col_scaler = ColumnTransformer([("num_data", StandardScaler(), numeric_cols),],
                               remainder='passthrough').
    set_output(transform="pandas")
df_cust_scaled = col_scaler.fit_transform(df_cust_trans)

```

StandardScaler() -transforms numeric features to have mean=0 and variance=1. This ensures all columns are on the same scale, preventing large-range features from dominating model training

```
[ ]: df_cust_scaled.describe()

[ ]:      num_data__avg_imputer__Session_Duration_Avg \
count                  5.000000e+04
mean                 -4.017409e-16
std                   1.000010e+00
min                  -2.036323e+00
25%                  -7.203269e-01
50%                  1.608806e-03
75%                  6.150224e-01
max                  2.647076e+00

      num_data__avg_imputer__Returns_Rate \
count                  5.000000e+04
mean                 4.033751e-16
std                   1.000010e+00
min                  -1.379414e+00
25%                  -7.591560e-01
50%                  -1.167458e-01
75%                  4.592082e-01
max                  3.338978e+00

      num_data__avg_imputer__Email_Open_Rate \
count                  5.000000e+04
mean                 1.818989e-17
std                   1.000010e+00
min                  -1.521137e+00
25%                  -7.563769e-01
50%                  -6.183219e-03
75%                  6.420424e-01
max                  2.630420e+00

      num_data__avg_imputer__Discount_Usage_Rate \
count                  5.000000e+04
mean                 -1.826095e-17
std                   1.000010e+00
min                  -1.826400e+00
25%                  -7.573748e-01
50%                  2.588236e-03
75%                  6.676646e-01
max                  2.503690e+00

      num_data__avg_imputer__Payment_Method_Diversity \
count                  5.000000e+04
mean                 -2.642508e-16
std                   1.000010e+00
min                  -1.251392e+00
```

```

25%           -3.270872e-01
50%           -3.270872e-01
75%           5.972177e-01
max            2.445828e+00

    num_data__avg_imputer__Credit_Balance \
count          5.000000e+04
mean           1.989520e-17
std             1.000010e+00
min            -1.716946e+00
25%            -6.982743e-01
50%            3.798246e-03
75%            6.144471e-01
max            2.662301e+00

    num_data__avg_imputer__Mobile_App_Usage  num_data__med_imputer__Age \
count          5.000000e+04           5.000000e+04
mean           1.830358e-16          -2.596323e-16
std             1.000010e+00          1.000010e+00
min            -2.113401e+00          -1.782070e+00
25%            -6.964337e-01          -6.985067e-01
50%            3.163578e-03           2.386910e-02
75%            5.958406e-01           6.559479e-01
max            2.670281e+00           2.461887e+00

    num_data__med_imputer__Wishlist_Items \
count          5.000000e+04
mean           1.251976e-16
std             1.000010e+00
min            -1.414477e+00
25%            -7.506744e-01
50%            -8.687184e-02
75%            5.769307e-01
max            3.232141e+00

    num_data__med_imputer__Customer_Service_Calls \
count          5.000000e+04
mean           -4.938272e-17
std             1.000010e+00
min            -1.781862e+00
25%            -6.396575e-01
50%            -2.589226e-01
75%            5.025473e-01
max            2.786957e+00

    num_data__mode_imputer__Pages_Per_Session \
count          5.000000e+04

```

```

mean                                5.641709e-16
std                                 1.000010e+00
min                                -2.042588e+00
25%                                 -6.891226e-01
50%                                 -1.366878e-01
75%                                 6.367208e-01
max                                2.625486e+00

    num_data__mode_imputer__Days_Since_Last_Purchase \
count                               5.000000e+04
mean                                7.929657e-17
std                                 1.000010e+00
min                                -9.840524e-01
25%                                 -7.360205e-01
50%                                 -3.108229e-01
75%                                 3.978396e-01
max                                3.834853e+00

    num_data__mode_imputer__Product_Reviews_Written \
count                               5.000000e+04
mean                                -4.131806e-17
std                                 1.000010e+00
min                                -1.205141e+00
25%                                 -7.607444e-01
50%                                 -3.163480e-01
75%                                 5.724447e-01
max                                3.238823e+00

    num_data__mode_imputer__Social_Media_Engagement_Score \
count                               5.000000e+04
mean                                8.697043e-17
std                                 1.000010e+00
min                                -1.208506e+00
25%                                 -9.411857e-01
50%                                 -9.701635e-02
75%                                 7.002547e-01
max                                2.623085e+00

    num_data__remainder__Membership_Years \
count                               5.000000e+04
mean                                -2.651035e-16
std                                 1.000010e+00
min                                -1.352749e+00
25%                                 -7.697732e-01
50%                                 -2.353785e-01
75%                                 4.933415e-01
max                                3.408222e+00

```

```

        num_data__remainder__Login_Frequency \
count                  5.000000e+04
mean                   3.478107e-17
std                    1.000010e+00
min                   -1.500666e+00
25%                   -7.242240e-01
50%                   -7.718872e-02
75%                   6.992536e-01
max                   2.640359e+00

        num_data__remainder__Cart_Abandonment_Rate \
count                  5.000000e+04
mean                   7.673862e-18
std                    1.000010e+00
min                   -2.537946e+00
25%                   -6.681313e-01
50%                   6.351092e-02
75%                   7.263663e-01
max                   2.108357e+00

        num_data__remainder__Total_Purchases \
count                  5.000000e+04
mean                   2.316369e-16
std                    1.000010e+00
min                   -1.771100e+00
25%                   -8.045689e-01
50%                   -1.602148e-01
75%                   6.452278e-01
max                   2.932685e+00

        num_data__remainder__Average_Order_Value \
count                  5.000000e+04
mean                   9.407586e-17
std                    1.000010e+00
min                   -1.704752e+00
25%                   -7.419689e-01
50%                   -1.443825e-01
75%                   5.811593e-01
max                   3.067877e+00

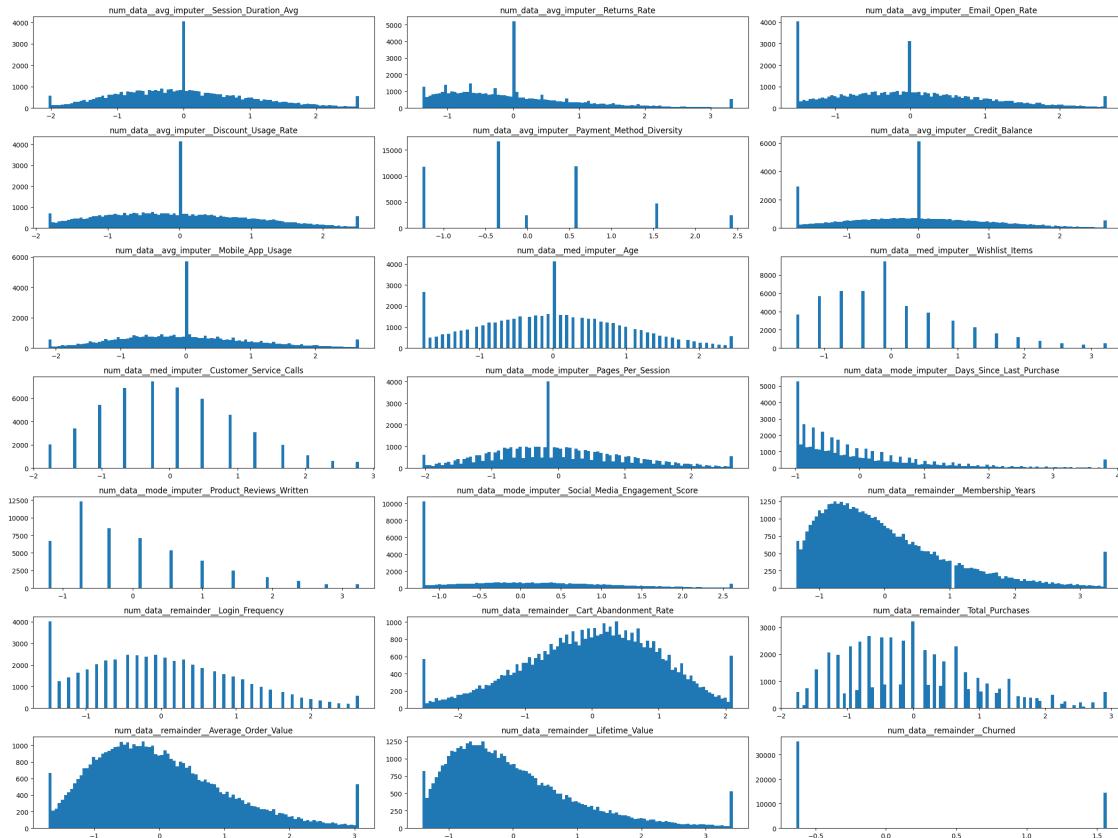
        num_data__remainder__Lifetime_Value  num_data__remainder__Churned
count                  5.000000e+04                  5.000000e+04
mean                   -9.762857e-17                 6.870948e-17
std                    1.000010e+00                 1.000010e+00
min                   -1.411868e+00                 -6.375498e-01
25%                   -7.370380e-01                 -6.375498e-01

```

50%	-2.173839e-01	-6.375498e-01
75%	5.050320e-01	1.568505e+00
max	3.422027e+00	1.568505e+00

Now, we'll re-plot the histogram plots after normalizing to examine the cleaned data.

```
[ ]: numeric_cols_scaled=df_cust_scaled.columns[df_cust_scaled.dtypes != 'object']
fig, axes = plt.subplots(nrows=7,ncols=3, figsize=(24,18))
axes=axes.ravel()
idx=0
for col in numeric_cols_scaled:
    if idx < len(axes): # Added a check to prevent IndexError
        axes[idx].set_title(col)
        axes[idx].hist(df_cust_scaled[col],bins=100)
    idx+=1
plt.tight_layout()
plt.show()
```



Now let's check how the numeric columns are related to each other using correlation.

Correlation measures how strongly two variables move together. A positive correlation means both increase or decrease together, while a negative correlation means one increases as the other

decreases.

```
[ ]: import seaborn as sns

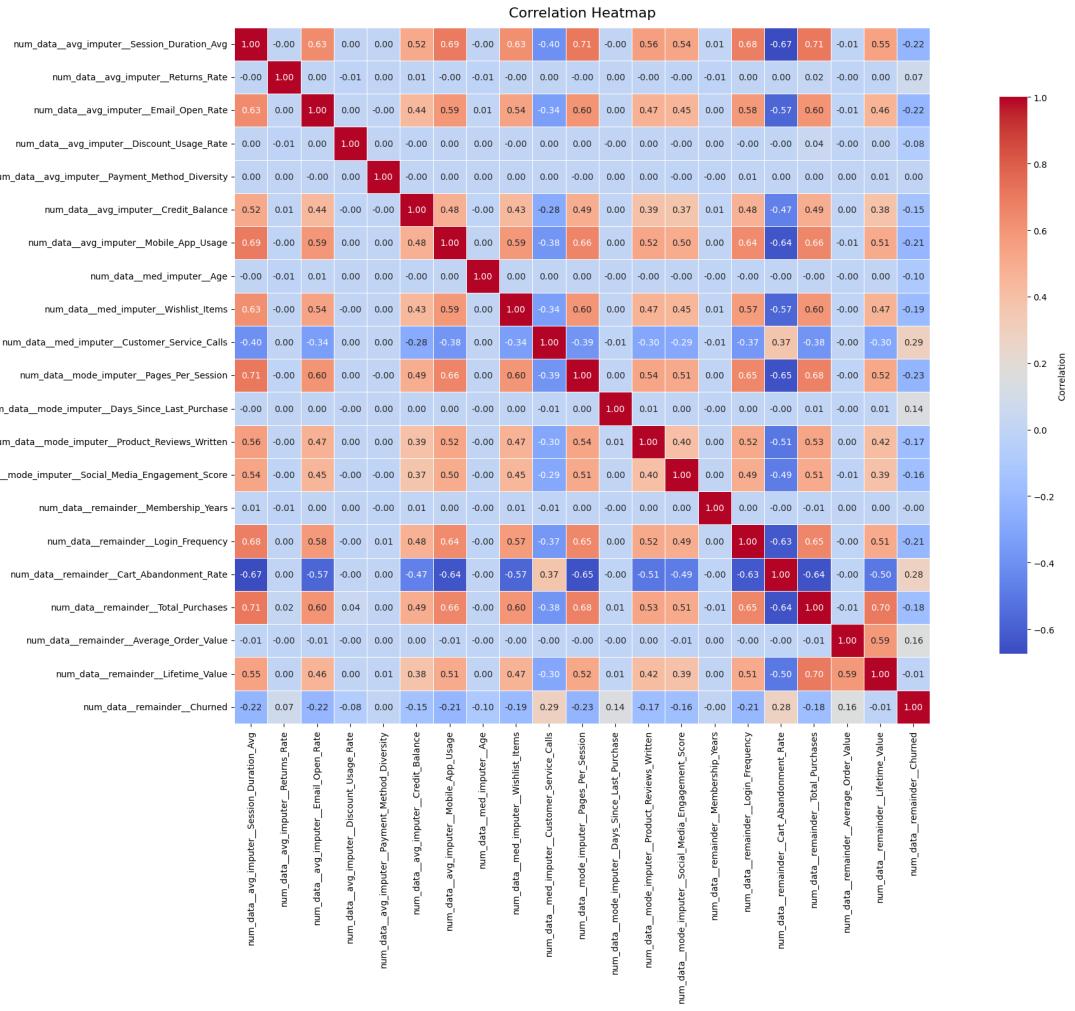
correlations = {}

for col in numeric_cols_scaled:
    correlations[col] = df_cust_scaled[numeric_cols_scaled].corr()[col].
    ↪to_dict()

# Convert dict-of-dicts to a DataFrame (matrix)
corr_df = pd.DataFrame(correlations)

# Ensure consistent ordering of rows/columns
corr_df = corr_df.loc[numeric_cols_scaled, numeric_cols_scaled]

# Plot heatmap
plt.figure(figsize=(24, 16))
sns.heatmap(corr_df,
            annot=True,
            cmap='coolwarm',
            fmt='.2f',
            linewidths=0.5,
            linecolor='white',
            square=True,
            cbar_kws={'shrink': .8, 'label': 'Correlation'})
plt.title('Correlation Heatmap', fontsize=16, pad=12)
plt.tight_layout()
plt.show()
```



Seaborn is a Python library that makes it easy to create beautiful charts. It's built on top of Matplotlib and adds simple functions for common plots like histograms, scatter plots, and heatmaps.

Next, we'll examine the non-numeric features in the dataset.

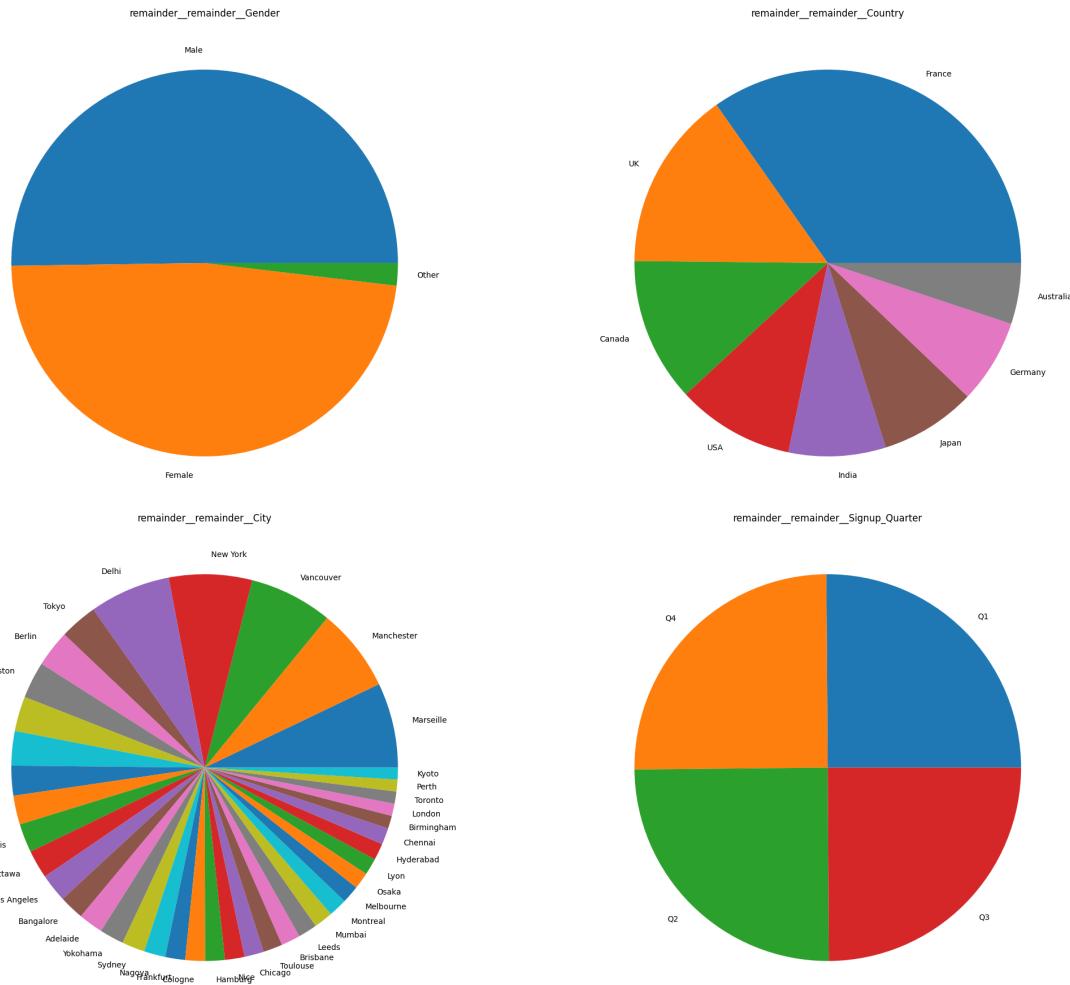
We will use pie chart to visualize categorical data. Each slice represents a category, and its size shows the proportion or frequency of that category. This makes it easy to compare groups at a glance and understand how the dataset is distributed across categories.

```
[ ]: categ_cols=df_cust_scaled.select_dtypes(include=['object', 'category']).columns.tolist()
fig, axes = plt.subplots(nrows=2,ncols=2, figsize=(24,18))
axes=axes.ravel()
idx=0
for col in categ_cols:
    axes[idx].set_title(col)
```

```

        axes[idx].pie(df_cust_scaled[col].value_counts(), labels=df_cust_scaled[col].
        ↪unique())
        idx+=1
plt.tight_layout()
plt.show()

```



Since machine learning models cannot directly process text data, we will convert categorical variables into numerical form using OneHotEncoder.

```
[ ]: !pip install category_encoders
```

```

Collecting category_encoders
  Downloading category_encoders-2.9.0-py3-none-any.whl.metadata (7.9 kB)
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (2.0.2)
Requirement already satisfied: pandas>=1.0.5 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (2.2.2)

```

```

Requirement already satisfied: patsy>=0.5.1 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (1.0.2)
Requirement already satisfied: scikit-learn>=1.6.0 in
/usr/local/lib/python3.12/dist-packages (from category_encoders) (1.6.1)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (1.16.3)
Requirement already satisfied: statsmodels>=0.9.0 in
/usr/local/lib/python3.12/dist-packages (from category_encoders) (0.14.6)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.12/dist-packages (from pandas>=1.0.5->category_encoders)
(2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-
packages (from pandas>=1.0.5->category_encoders) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-
packages (from pandas>=1.0.5->category_encoders) (2025.3)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-
packages (from scikit-learn>=1.6.0->category_encoders) (1.5.3)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-
learn>=1.6.0->category_encoders) (3.6.0)
Requirement already satisfied: packaging>=21.3 in
/usr/local/lib/python3.12/dist-packages (from
statsmodels>=0.9.0->category_encoders) (25.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-
packages (from python-dateutil>=2.8.2->pandas>=1.0.5->category_encoders)
(1.17.0)
Downloading category_encoders-2.9.0-py3-none-any.whl (85 kB)
  85.9/85.9 kB
  6.5 MB/s eta 0:00:00
Installing collected packages: category_encoders
Successfully installed category_encoders-2.9.0

```

```
[ ]: from category_encoders import OneHotEncoder

col_trans = ColumnTransformer([('categ_data', OneHotEncoder(), categ_cols)],
                             remainder='passthrough').
    ↪set_output(transform="pandas")
df_cust_categ = col_trans.fit_transform(df_cust_scaled)
```

Let's check what happens to our categorical columns once we use OneHotEncoder.

```
[ ]: df_cust_categ.info()
```

<class 'pandas.core.frame.DataFrame'>	
RangeIndex: 50000 entries, 0 to 49999	
Data columns (total 76 columns):	
# Column	Non-Null
Count Dtype	

```

--- -----
-----  -----
0  categ_data__remainder__remainder__Gender_1          50000
non-null  int64
1  categ_data__remainder__remainder__Gender_2          50000
non-null  int64
2  categ_data__remainder__remainder__Gender_3          50000
non-null  int64
3  categ_data__remainder__remainder__Country_1         50000
non-null  int64
4  categ_data__remainder__remainder__Country_2         50000
non-null  int64
5  categ_data__remainder__remainder__Country_3         50000
non-null  int64
6  categ_data__remainder__remainder__Country_4         50000
non-null  int64
7  categ_data__remainder__remainder__Country_5         50000
non-null  int64
8  categ_data__remainder__remainder__Country_6         50000
non-null  int64
9  categ_data__remainder__remainder__Country_7         50000
non-null  int64
10 categ_data__remainder__remainder__Country_8         50000
non-null  int64
11 categ_data__remainder__remainder__City_1           50000
non-null  int64
12 categ_data__remainder__remainder__City_2           50000
non-null  int64
13 categ_data__remainder__remainder__City_3           50000
non-null  int64
14 categ_data__remainder__remainder__City_4           50000
non-null  int64
15 categ_data__remainder__remainder__City_5           50000
non-null  int64
16 categ_data__remainder__remainder__City_6           50000
non-null  int64
17 categ_data__remainder__remainder__City_7           50000
non-null  int64
18 categ_data__remainder__remainder__City_8           50000
non-null  int64
19 categ_data__remainder__remainder__City_9           50000
non-null  int64
20 categ_data__remainder__remainder__City_10          50000
non-null  int64
21 categ_data__remainder__remainder__City_11          50000
non-null  int64
22 categ_data__remainder__remainder__City_12          50000
non-null  int64

```

23	categ_data__remainder__remainder__City_13	50000
	non-null int64	
24	categ_data__remainder__remainder__City_14	50000
	non-null int64	
25	categ_data__remainder__remainder__City_15	50000
	non-null int64	
26	categ_data__remainder__remainder__City_16	50000
	non-null int64	
27	categ_data__remainder__remainder__City_17	50000
	non-null int64	
28	categ_data__remainder__remainder__City_18	50000
	non-null int64	
29	categ_data__remainder__remainder__City_19	50000
	non-null int64	
30	categ_data__remainder__remainder__City_20	50000
	non-null int64	
31	categ_data__remainder__remainder__City_21	50000
	non-null int64	
32	categ_data__remainder__remainder__City_22	50000
	non-null int64	
33	categ_data__remainder__remainder__City_23	50000
	non-null int64	
34	categ_data__remainder__remainder__City_24	50000
	non-null int64	
35	categ_data__remainder__remainder__City_25	50000
	non-null int64	
36	categ_data__remainder__remainder__City_26	50000
	non-null int64	
37	categ_data__remainder__remainder__City_27	50000
	non-null int64	
38	categ_data__remainder__remainder__City_28	50000
	non-null int64	
39	categ_data__remainder__remainder__City_29	50000
	non-null int64	
40	categ_data__remainder__remainder__City_30	50000
	non-null int64	
41	categ_data__remainder__remainder__City_31	50000
	non-null int64	
42	categ_data__remainder__remainder__City_32	50000
	non-null int64	
43	categ_data__remainder__remainder__City_33	50000
	non-null int64	
44	categ_data__remainder__remainder__City_34	50000
	non-null int64	
45	categ_data__remainder__remainder__City_35	50000
	non-null int64	
46	categ_data__remainder__remainder__City_36	50000
	non-null int64	

47 categ_data__remainder__remainder__City_37	50000
non-null int64	
48 categ_data__remainder__remainder__City_38	50000
non-null int64	
49 categ_data__remainder__remainder__City_39	50000
non-null int64	
50 categ_data__remainder__remainder__City_40	50000
non-null int64	
51 categ_data__remainder__remainder__Signup_Quarter_1	50000
non-null int64	
52 categ_data__remainder__remainder__Signup_Quarter_2	50000
non-null int64	
53 categ_data__remainder__remainder__Signup_Quarter_3	50000
non-null int64	
54 categ_data__remainder__remainder__Signup_Quarter_4	50000
non-null int64	
55 remainder__num_data__avg_imputer__Session_Duration_Avg	50000
non-null float64	
56 remainder__num_data__avg_imputer__Returns_Rate	50000
non-null float64	
57 remainder__num_data__avg_imputer__Email_Open_Rate	50000
non-null float64	
58 remainder__num_data__avg_imputer__Discount_Usage_Rate	50000
non-null float64	
59 remainder__num_data__avg_imputer__Payment_Method_Diversity	50000
non-null float64	
60 remainder__num_data__avg_imputer__Credit_Balance	50000
non-null float64	
61 remainder__num_data__avg_imputer__Mobile_App_Usage	50000
non-null float64	
62 remainder__num_data__med_imputer__Age	50000
non-null float64	
63 remainder__num_data__med_imputer__Wishlist_Items	50000
non-null float64	
64 remainder__num_data__med_imputer__Customer_Service_Calls	50000
non-null float64	
65 remainder__num_data__mode_imputer__Pages_Per_Session	50000
non-null float64	
66 remainder__num_data__mode_imputer__Days_Since_Last_Purchase	50000
non-null float64	
67 remainder__num_data__mode_imputer__Product_Reviews_Written	50000
non-null float64	
68 remainder__num_data__mode_imputer__Social_Media_Engagement_Score	50000
non-null float64	
69 remainder__num_data__remainder__Membership_Years	50000
non-null float64	
70 remainder__num_data__remainder__Login_Frequency	50000
non-null float64	

```

71 remainder__num_data__remainder__Cart_Abandonment_Rate      50000
non-null float64
72 remainder__num_data__remainder__Total_Purchases          50000
non-null float64
73 remainder__num_data__remainder__Average_Order_Value    50000
non-null float64
74 remainder__num_data__remainder__Lifetime_Value         50000
non-null float64
75 remainder__num_data__remainder__Churned                 50000
non-null float64
dtypes: float64(21), int64(55)
memory usage: 29.0 MB

```

As we can see each value in category features has converted into a column (3 value in Gender is completed into 3 columns)

```
[ ]: print(df_cust_categ[['categ_data__remainder__remainder__Gender_1',
                           'categ_data__remainder__remainder__Gender_2',
                           'categ_data__remainder__remainder__Gender_3']])
                           ↵head(5).to_string())
```

	categ_data__remainder__remainder__Gender_1	categ_data__remainder__remainder__Gender_2	categ_data__remainder__remainder__Gender_3
0	1	0	0
0	0	1	0
1	0	0	1
0	0	0	1
2	0	1	0
1	0	0	1
3	0	0	1
1	0	1	0
4	1	0	0
0	0	0	1

Apart from above, listed below is some encoder used for non-numerical data

- Label Encoding: Assigns each category a unique integer. Best for ordered data where numbers reflect ranking.
- One-Hot Encoding: Creates binary columns for each category. Useful for nominal data without any natural order.
- Ordinal Encoding: Maps categories to integers based on order. Ideal when categories have a clear progression.
- Binary Encoding: Converts categories into binary digits across fewer columns. Helps reduce dimensionality for large sets.
- Target/Mean Encoding: Replaces categories with target variable averages. Effective but must avoid data leakage.
- Frequency/Count Encoding: Substitutes categories with their occurrence count. Simple and works well with tree-based models.

Note: These are some of the ways to pre-process data before using the same for training the model.

#End of Chapter-2